

Beginning Steps in Network Data Analysis: Understanding Crossposting between Reddit Communities

Carlyle Morgan

Introduction

The Stanford Network Analysis Project(SNAP) hosts a dataset that contains information about reddit ‘crossposting’, in which a post on one subcommunity(or ‘Subreddit’) of the site posts a hyperlink to another subcommunity of the site. According to SNAP, the Reddit data is formatted such that “each hyperlink is annotated with three properties: the timestamp, the sentiment of the source community post towards the target community post, and the text property vector of the source post. The network is directed, signed, temporal, and attributed”. From these attributes, one can understand

Cleaning

```
library(igraph)
```

```
## Warning: package 'igraph' was built under R version 4.0.5
```

```
##
```

```
## Attaching package: 'igraph'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      decompose, spectrum
```

```
## The following object is masked from 'package:base':
```

```
##
```

```
##      union
```

```
library(dplyr)
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:igraph':
```

```
##
```

```
##      as_data_frame, groups, union
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      intersect, setdiff, setequal, union
```

```
library(stringr)
```

```
rh<-read.csv("C:/Users/sc_mo/Documents/R/Projects/Stats Research/reddithyperlinks.tsv", sep = "\t")
```

```
g.rh<- graph.data.frame(rh)
is.simple(g.rh)
```

```
## [1] FALSE
```

```
##Weight according to the number of connected edges
```

```
g.rh<-as.undirected(g.rh, mode = 'collapse')
E(g.rh)$weight<-1
g.rh.weighted<-simplify(g.rh)
```

Exploratory Data Analysis

Overall Edge and Vertex Summary Statistics

Having simplified and applying weighting to our network, it might be helpful to explore the relationship between the degrees and strength of our vertices.

Which subreddits have the greatest number of instances of starting and receiving crossposts? This is concerned with how we decided to weight each edge, and thus nodes with greater summed weights have

```
head(sort(graph.strength(g.rh.weighted), decreasing=TRUE),25)
```


```
##      askreddit      todayilearned      writingprompts      iama
##      3079          2371          2242          2093
##      subredditdrama      outoftheloop      pics      leagueoflegends
##      1887          1232          1196          1117
##      adviceanimals      tipofmypenis      copypasta      videos
##      1117          1096          1031          959
##      wtf          funny      explainlikeimfive      mhoc
##      924          897          863          857
##      circlebroke      worldnews      conspiracy      gaming
##      807          806          779          763
##      drama      hailcorporate      circlejerkcopypasta      shitliberalssay
##      709          701          700          698
##      dogecoin
##      638
```

```
head(sort(degree(g.rh.weighted), decreasing=TRUE),25)
```

```
##      askreddit      iama      subredditdrama      writingprompts
##      2336          1839          1598          1043
##      outoftheloop      pics      videos      todayilearned
##      986          956          881          816
##      funny      gaming      leagueoflegends      copypasta
##      783          719          708          669
##      conspiracy      worldnews      drama      explainlikeimfive
##      666          665          642          617
##      mhoc      legaladvice      news      subredditoftheday
##      607          590          587          573
##      the_donald      pcmasterrace      bitcoin      tipofmypenis
##      565          500          487          485
##      dogecoin
##      484
```


```
hist(graph.strength(g.rh.weighted), col="red", xlab="Vertex Strength",
     ylab="Frequency",
     xlim = c(0,200), breaks = seq(0,10000,by=10))
```

Histogram of graph.strength(g.rh.weighted)




```
hist(degree(g.rh.weighted), col="blue", xlab="Vertex Degree",  
     ylab="Frequency", xlim = c(0,200), breaks = seq(0,10000,by=10))
```

Histogram of degree(g.rh.weighted)



```
d.rh <- degree(g.rh.weighted)
dd.rh <- degree.distribution(g.rh.weighted)
d <- 1:max(d.rh)-1
ind <- (dd.rh != 0)
plot(d[ind], dd.rh[ind], log="xy", col="purple", xlab=c("Log-Degree"), ylab=c("Log-Intensity"), main="L
```

Log-Log Degree Distribution



Component Analysis

```
is.connected(g.rh.weighted)
```

```
## [1] FALSE
```

```
parts<-decompose.graph(g.rh.weighted)
table(sapply(parts, vcount))
```

```
##
##      2      3      4      5      8      9     42 34671
## 450    33     3     7     1     1     1     1
```

There is clearly the presence of a giant component in this data.

```
vcount(g.rh.weighted)
```

```
## [1] 35776
```

```
g.rh.weighted<-decompose.graph(g.rh.weighted)[[1]]
vcount(g.rh.weighted)
```

```
## [1] 34671
```

```
is.connected(g.rh.weighted)
```

```
## [1] TRUE
```

```
g.rh.weighted<-delete_vertices(g.rh.weighted, degree(g.rh.weighted) <= 100)
```

```
average.path.length(g.rh.weighted)
```

```
## [1] 1.80745
```

With such a small average path length,

Hierarchical Clustering and Graph Partitioning

```
greedy_rh<-fastgreedy.community(g.rh.weighted)  
sizes(greedy_rh)
```

```
## Community sizes
```

```
##  1  2  3  4  5  6  7
```

```
## 10  9 137 13 37 142  2
```

```
plot(greedy_rh,g.rh.weighted, asp = 0, vertex.size = 5, vertex.label.cex = 0.8, vertex.label=NA)
```

